



# *Aas Trailblazers*

## Unleashing insights

# Index Options

## Row Store

- Heap
- Clustered Index
- Non-Clustered Index

## Clustered Columnstore Index

- Default if not defined in CREATE TABLE

# Row store vs. Column store

Row store for B-Tree or Heap										
Row 1	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 2	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 3	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 4	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 5	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Page 1										
Row 6	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 7	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 8	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
.....	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row n	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Page 2										

Column Store Index										
Row 1	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 2	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 3	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 4	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 5	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 6	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 7	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 8	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
.....	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row n	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Page 1 Page 2 Page 3 Page 4 Page 5 Page 6 Page 7 Page 8 Page 9 Page 10										

# Row store: Heaps

- Just a collection of pages
  - Append only
  - No ordering
  - Minimal fragmentation
- Fast target for loading
- Much slower when used for small range scans & single selects

## Understanding Heaps

Deletes are not physically removed from a heap until it is rebuilt. As data is added SQL allocates more pages to the table.

# Row store: Indexes (clustered and non-clustered)

- Use a Balanced Tree (b-tree) to organize data
- Get fragmented over time
  - Require more maintenance
- Leaf level
  - Clustered index: data
  - Non-clustered index: Row ID (RID) – may require lookup of data
- Ideal for limited range scans & singleton selects
- Slower for table scans / partition scans / loading

# Columnstore: Index

Group rows into batches of ~1M rows

- This is called a row group

Compress/Encode each column

- Lower selectivity compresses more efficiently
- Primitive data types (int, date, short string, etc.) compresses efficiently
- Long string fields do NOT compress well

Typically, up to 10x  
compression and  
query performance  
up to 10x\*

\* <https://docs.microsoft.com/en-us/sql/relational-databases/indexes/columnstore-indexes-overview?view=sql-server-ver15>

# Columnstore: Design and storage

# Data

# Row Group

# Segments

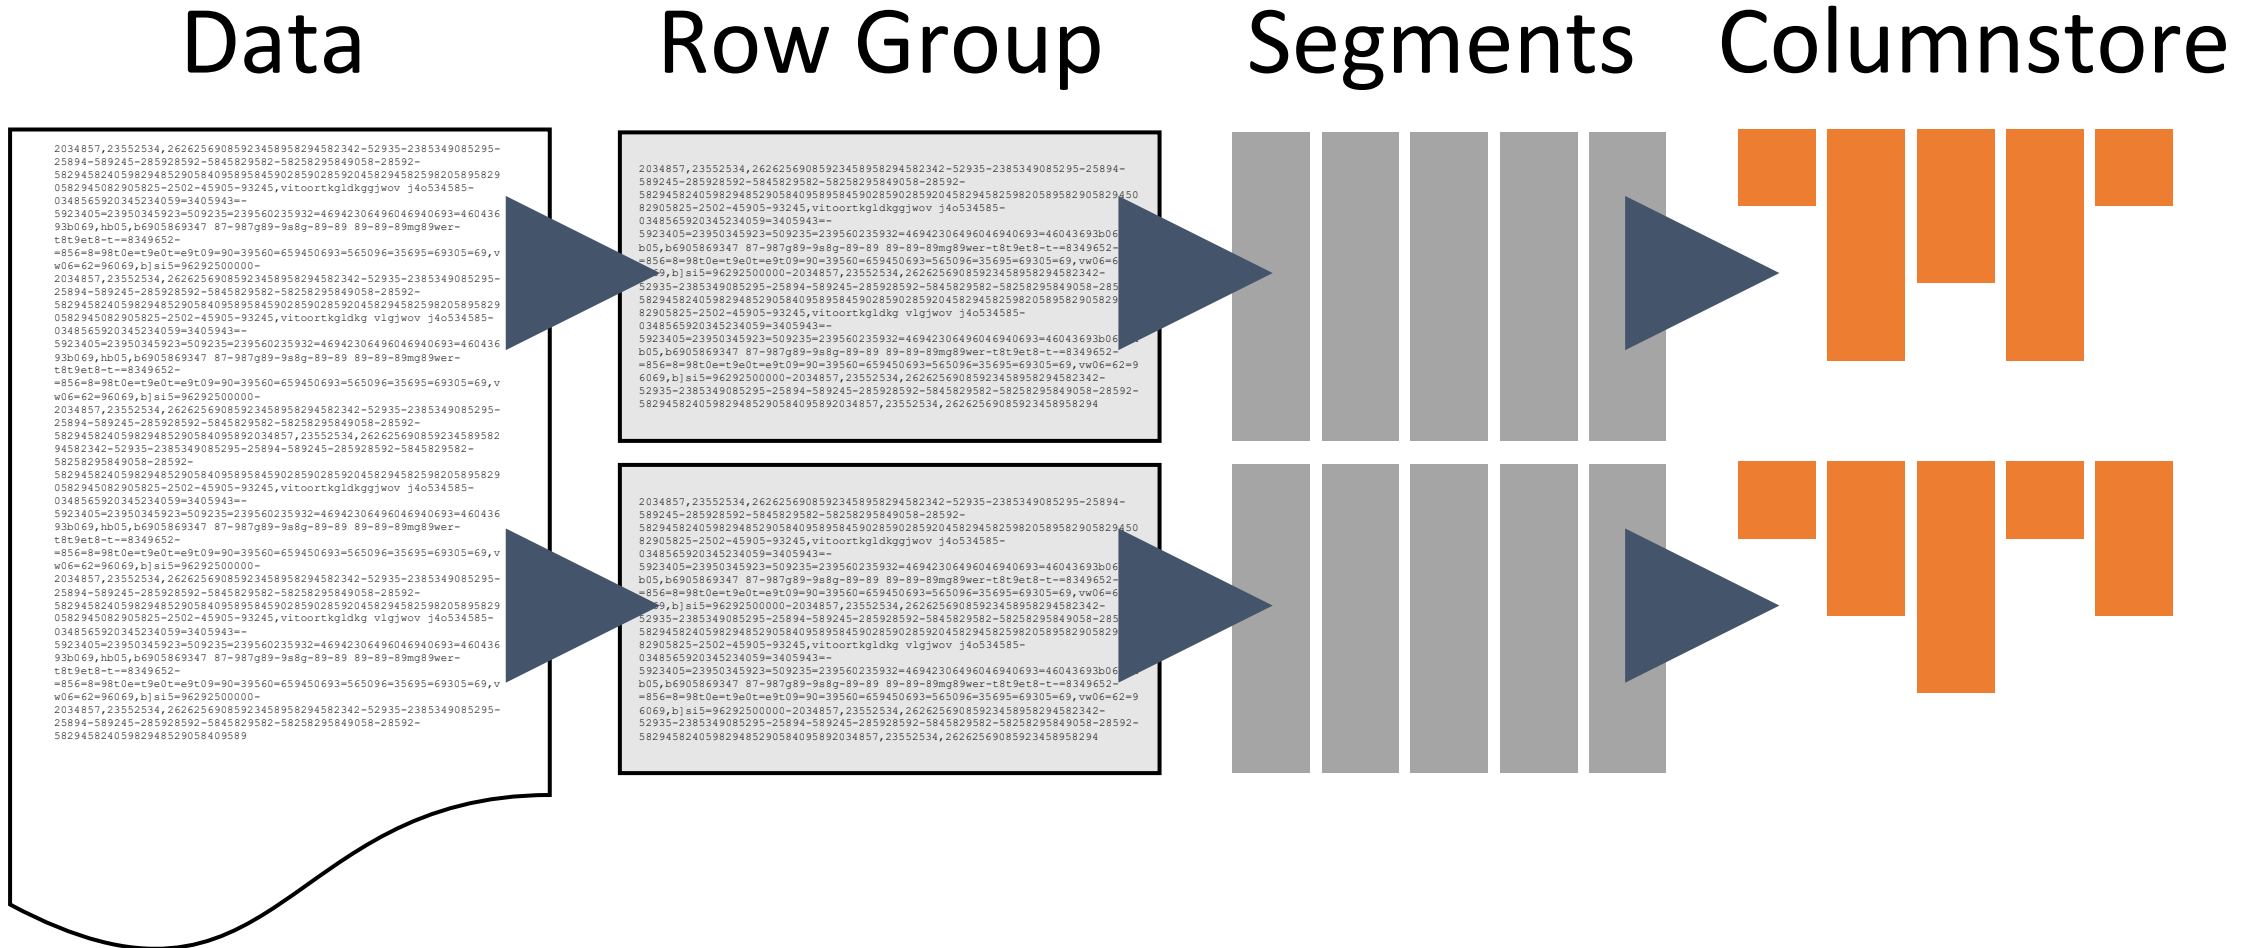
# Columnstore



Demo



# Columnstore: Design and storage



# Columnstore: Row Groups

Exist in one of these states

- **Open:** In delta store; accepting new rows, data is not compressed or indexed
- **Closed:** In delta store; not accepting new rows
- **Compressed:** In Column Store format
- **Tombstone:** a rowgroup with all the data deleted

System View: **sys.pdw\_nodes\_column\_store\_row\_groups**

# Columnstore: Column Segments

Example of a segment from `sys.pdw_nodes_column_store_segments`

segment_id	row_count	base_id	min_data_id	max_data_id
0	1048576	19568138	19568141	30426164
1	1048576	19757017	19757020	31243669
2	1048576	19788588	19788591	31341348
3	1048576	19778863	19778866	31349845
4	1048576	19574032	19574035	31365043
5	1048576	19773033	19773036	31330103

Demo

# Columnstore: Updating

CCIs perform **large block compression**

- Maximises compression
- Expensive to update

CCIs use two mechanisms to **handle updates**


- Delete bitmap
- Delta stores

# Columnstore: Delta Store

## Page compressed b-tree i.e. a row store

- B-tree on unique integer row ID
- Matches user columns defined in the CCI
- Aligned to underlying CCI partition
- Created only if needed

Max # rows  
delta store  
1,048,576



# Columnstore: Delta Store to Columnstore

Once a row group has closed it can be converted to column store format:

As an **online operation**

- Tuple Mover: a background process
- ALTER INDEX..REORGANIZE WITH (COMPRESS\_ALL\_ROW\_GROUPS = ON);

As an **offline operation**

- ALTER INDEX..REBUILD

# Columnstore: Small DML

## **Inserts**

- Small inserts are written first to the delta store
- Bulk-load based on some internal threshold writes as column store

## **Deletes**

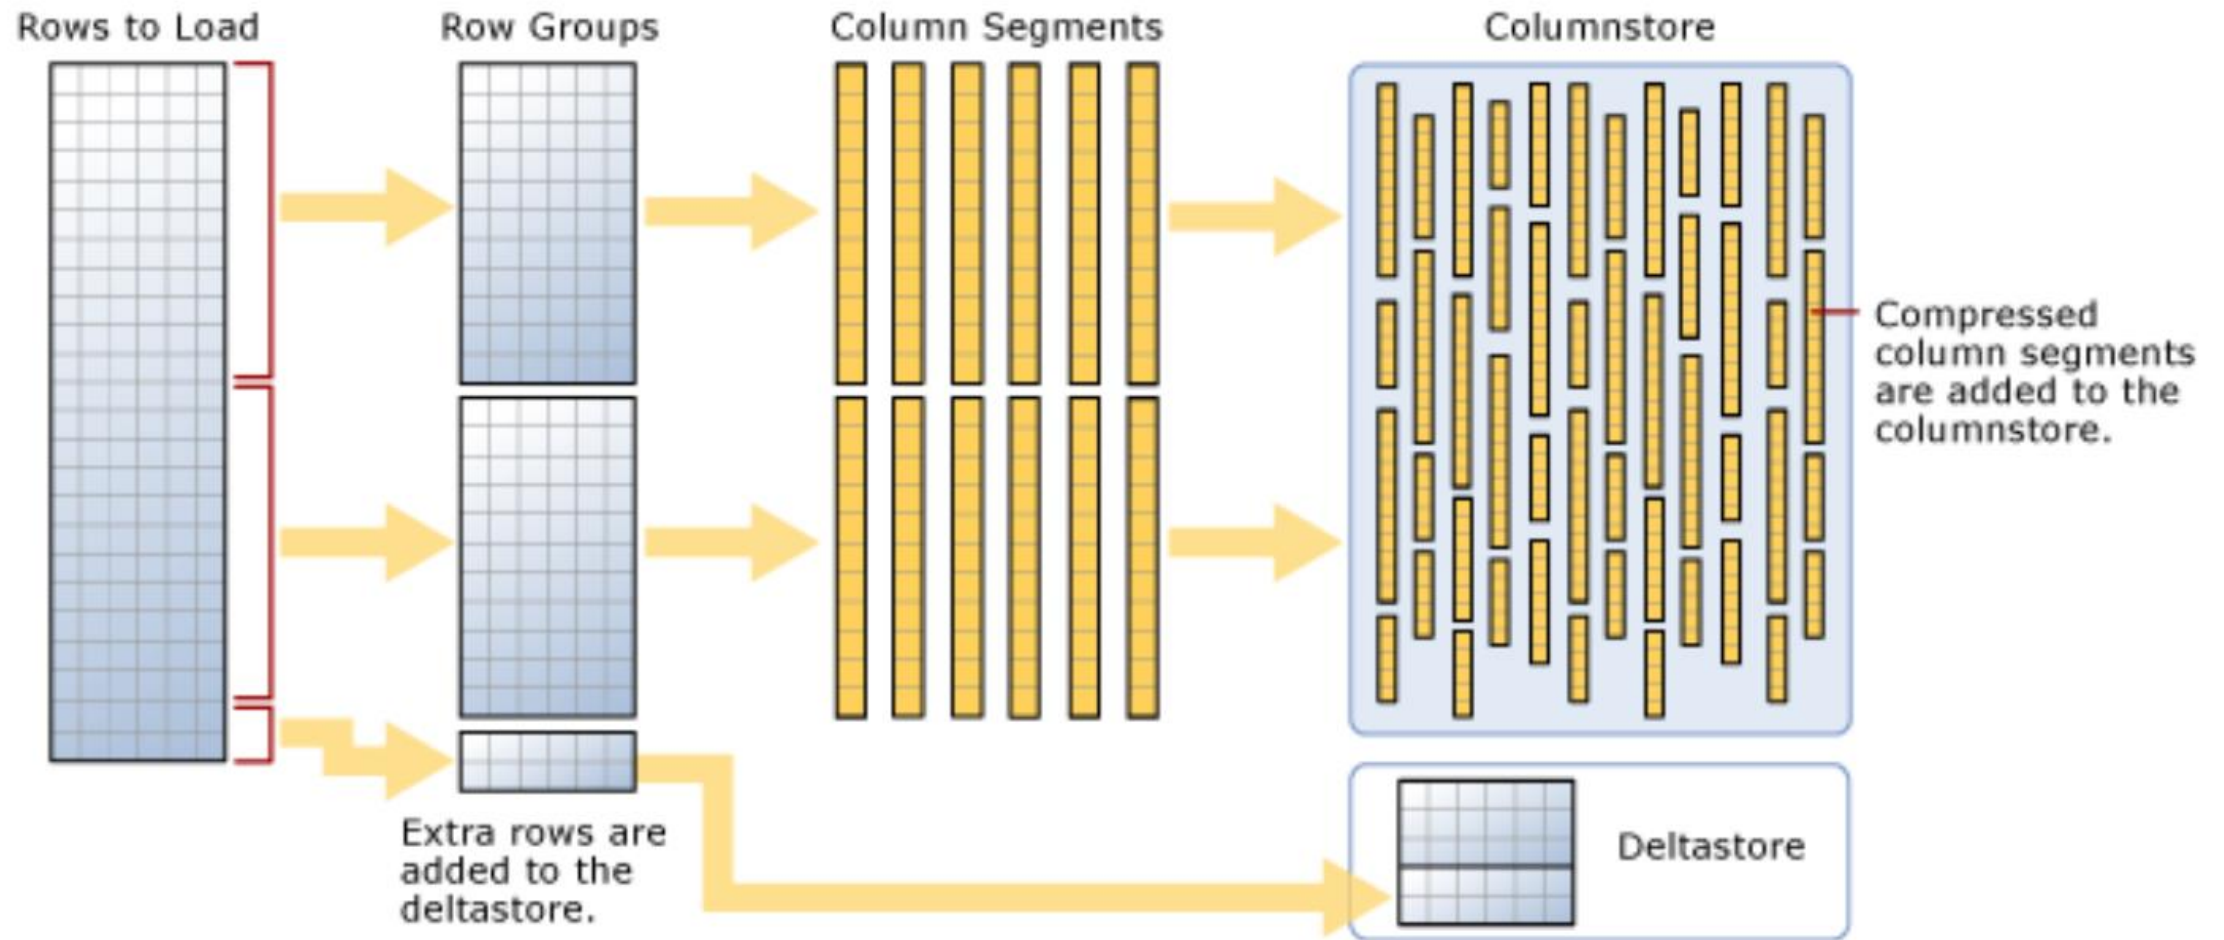
- Logical against rows in column store
- Physically against rows in the delta store

## **Updates**

- Converted to a logical delete and an insert



# Columnstore: Design and storage



Demo

# Optimizing with Indexes: Choose right type

## Clustered Columnstore (Default)

- Optimal choice for large tables
- Limits scans to columns in the query
- Optimal compression
- Slower to load than Heap
- Keep partitions large enough to compress (> 1 million rows)

## Heap

- Optimal choice for temporary or staging tables
- Fastest load performance

## Clustered Index

- Optimal for tables < 60M rows
- Sorting operation slows-down load

## Non-clustered Index

- **Use sparingly**
- Optimize single row lookups
- Will slow-down load